

[Return to Classroom](#)

Generate TV Scripts

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

Congratulations 🎉🎉🎉

- Your submission reveals that you have made an **excellent effort** in finishing this project, especially the batching, model architecture and hyperparameters.
- Very good hyperparameters and decreasing cross entropy loss. 👍
- Please go through the additional suggestions in the rubric below.
- I wish you all the best for next adventures 🚀

Few references to explore:

- [Colah's Blog](#) A visual explanation of LSTMs you want to look at to understand it more.
- [Andrej Karpathy](#) Andrej is director of AI and AutoPilot Vision at Tesla, this link explains the unreasonable effectiveness of RNN
- [Rohan Kapur](#) , This link is an overall explanation of RNNs that you may find useful and insightful.

(feel free to reach out to mentors in the [knowledge forum](#) regarding any question or confusion)

Keep up the good work 👏 Stay Udacious 🍷

All Required Files and Tests

✔	The project submission contains the project notebook, called "d1nd_tv_script_generation.ipynb".
	All required files are present 👍 Bonus tips: <ul style="list-style-type: none">• It is recommended to export your conda environment into environment.yml file using command conda env export -f environment.yml, so that you can recreate your conda environment later.• While submitting this to any version control system like Github, make sure to include helper, data and environment files and exclude and temp files. It will help you in future if you want to re-execute it. Some guideline for best practice.
✔	All the unit tests in project have passed.
	All the unit tests in project have passed. 👍 Donald Knuth (a famous computer science pioneer) once famously said about unit tests: <i>"Beware of bugs in the above code; I have only proved it correct, not tried it."</i> Article on unit-test in machine learning system here

Pre-processing Data

✔	The function <code>create_lookup_tables</code> create two dictionaries: <ul style="list-style-type: none">• Dictionary to go from the words to an id, we'll call <code>vocab_to_int</code>• Dictionary to go from the id to word, we'll call <code>int_to_vocab</code> The function <code>create_lookup_tables</code> return these dictionaries as a tuple (<code>vocab_to_int</code> , <code>int_to_vocab</code>).
	Good job! 🍷 <ul style="list-style-type: none">• The <code>Counter</code> is a convenient way to get the information needed for that approach, but it has some extra overhead we don't need, so you could just use a <code>set</code> instead: <code>set(text)</code>. The sorting is also unnecessary.• You also only need to enumerate once, if you create both dicts inside a single for loop. All of these things will save some compute power/time. Alternatively, it could be implemented like this: <pre>vocab = set(text) vocab_to_int, int_to_vocab = {}, {} for i, w in enumerate(vocab): vocab_to_int[w] = i int_to_vocab[i] = w return (vocab_to_int, int_to_vocab)</pre>
✔	The function <code>token_lookup</code> returns a dict that can correctly tokenizes the provided symbols.

Batching Data

✔	The function <code>batch_data</code> breaks up word id's into the appropriate sequence lengths, such that only complete sequence lengths are constructed.
	Good Job 🍷 <ul style="list-style-type: none">• The implementation breaking up word id's into the appropriate sequence lengths
✔	In the function <code>batch_data</code> , data is converted into Tensors and formatted with TensorDataset.
	This implementation is basically responsible for loading sequenced data into Tensors in order for PyTorch's TensorDataset utility to generate the dataset. <pre>data = TensorDataset(feature_tensors, target_tensors)</pre> Check this dataloading tutorial
✔	Finally, <code>batch_data</code> returns a DataLoader for the batched training data.
	The unit test output tensor verifies the implementation 👍 <ul style="list-style-type: none">• The function of DataLoader is to combine a dataset and a sampler, and finally provides an iterable over the given dataset. Feature like automatic batching are also supported. Check details of DataLoader here• Adding <code>shuffle=True</code> in the dataloader is allowing you to add randomness in the training sequences.

Build the RNN

✔	The RNN class has complete <code>__init__</code> , <code>forward</code> , and <code>init_hidden</code> functions.
	👍 <ul style="list-style-type: none">• <code>__init__</code> , <code>forward</code> and <code>init_hidden</code> functions are complete, a good model architecture• it is recommended to remove unnecessary/unused methods from the code. e.g. <code>dropout</code>• RNN implements an LSTM Layer, and initializes it appropriately. <pre>self.lstm = nn.LSTM(embedding_dim, hidden_dim, n_layers, dropout=dropout, batch_first=True)</pre> <ul style="list-style-type: none">• we can also use GRU as well in place of LSTM:<pre>self.embedding = nn.Embedding(self.vocab_size, self.embedding_dim) self.GRU = nn.GRU(self.embedding_dim, self.hidden_dim, self.n_layers, batch_first=True, dropout=self.dropout) self.fc = nn.Linear(self.hidden_dim, self.output_size)</pre>• Check this article: Difference between LSTM and GRU
✔	The RNN must include an LSTM or GRU and at least one fully-connected layer. The LSTM/GRU should be correctly initialized, where relevant.
	The ideal structure is as follows: <ul style="list-style-type: none">• Embedding layer (nn.Embedding) before the LSTM or GRU layer.• The fully-connected layer comes at the end to get our desired number of outputs.• It is also recommended to not use a dropout after LSTM and before FC layer, as the drop out is already incorporated in the LSTMs, A lot of students adds it and then end up finding convergence difficult. The added layer could cause the model to lose key information that's needed to improve the model's performance.

RNN Training

✔	<ul style="list-style-type: none">• Enough epochs to get near a minimum in the training loss, no real upper limit on this. Just need to make sure the training loss is low and not improving much with more training.• Batch size is large enough to train efficiently, but small enough to fit the data in memory. No real "best" value here, depends on GPU memory usually.• Embedding dimension, significantly smaller than the size of the vocabulary, if you choose to use word embeddings• Hidden dimension (number of units in the hidden layers of the RNN) is large enough to fit the data well. Again, no real "best" value.• n_layers (number of layers in a GRU/LSTM) is between 1-3.• The sequence length (seq_length) here should be about the size of the length of sentences you want to look at before you generate the next word.• The learning rate shouldn't be too large because the training algorithm won't converge. But needs to be large enough that training doesn't take forever.
	🚀 <ul style="list-style-type: none">• Enough epochs to get near a minimum in the training loss.• Batch size is large enough to train efficiently. In order to use the GPU more efficiently, we can always try to set a value that is a power of two (e.g. 64 or 128 or 256)• Sequence length is about the size of the lengths of sentences we want to generate. Considering the fact that there are approximately an average of 11.504 words per line and 15.248 sentences in each scene• Size of embedding is in the range of [200-300]. The vocab contained ~46,367 unique words. Now you can try to cut this down significantly by 98% to 1000 embeddings. For example, Google's news word vectors, the GloVe vectors, and other word vectors are usually in the range 50 to 300• Learning rate seems good based on other hyper parameter• Hidden Dimension: 128-256 hidden dimensions to give the network a solid amount of features/states to learn from. recommendation on how to select Hidden Dimension <code>hidden_dim = 5*int(len(train_loader.dataset) / (embedding_dim + output_size))</code>• Number of layers is in between 1-3 as suggested in the project. <p>Your efforts shows that you have really have thought about it to get an optimized value 🔥</p>

✔	The printed loss should decrease during training. The loss should reach a value lower than 3.5.
	🔗 excellent decreasing loss .. <pre>Epoch: 14/15 Loss: 3.498869930744171 Epoch: 14/15 Loss: 3.4967117638587952 Epoch: 14/15 Loss: 3.5091017441749575 Epoch: 15/15 Loss: 3.476343840145351 Epoch: 15/15 Loss: 3.4058063492774964</pre>
✔	There is a provided answer that justifies choices about model size, sequence length, and other parameters.
	good answer, I suggest improving it by justifying each of the choices in the hyperparameters. <p>The act of elaborating your approach often leads to a deeper understanding of the material 🧐</p>

Generate TV Script

✔	The generated script can vary in length, and should look structurally similar to the TV script in the dataset.
	It doesn't have to be grammatically correct or make sense.
	well generated fun script! 🍷 <ul style="list-style-type: none">• all the lines are making sense• sentences are grammatically intact

📄 DOWNLOAD PROJECT

Rate this review

START

RETURN TO PATH